

Carla Prieto

AI & Machine Learning Engineer | Founder, Yoseph AI

carla.prieto@berkeley.edu | 206.581.4834 | linkedin.com/in/carlaprieto | github.com/Carla08 | carlaprieto.com | San Francisco Bay Area, CA

EDUCATION

Master of Science in Entrepreneurship and Innovation (MSEI), University of San Francisco

In Progress

Master of Information and Data Science, University of California, Berkeley

August 2024

Bachelor of Science in Information Technologies Engineering, Tec de Monterrey (ITESM)

December 2016

LANGUAGES AND TECHNOLOGIES

Proficient: Python, TensorFlow, Docker, Kubernetes, Airflow, NumPy, SciPy, Scikit-learn, Pandas, AWS (SageMaker, Bedrock, EC2, S3, Lambda), Google Cloud (GCP), Azure, Java, Kafka, SQL, Jupyter, Git, PyCharm, OpenAI API.

Exposure: PyTorch, Hugging Face Transformers, TensorBoard, AWS Glue, Weights & Biases, ML Flow

Machine Learning: LLMs, RAG, NLP, Embeddings, BERT, Distributed Training (DSP, FSDP), Supervised Learning, Deep Learning, CNNs, Classification, Clustering, Feature Engineering, Dimensionality Reduction, Cross-Validation, Model Evaluation

PROFESSIONAL EXPERIENCE

Subbase | AI Engineer | Seattle, WA

Dec 2025 – Now

- Led development of a new LLM-powered extraction system for invoices and quotes, improving structured data accuracy **from ~50% to 85%** and significantly reducing manual corrections.
- Replaced legacy extraction tooling, generating **\$40K+ in annual cost savings** while improving system reliability and scalability.
- Designed and implemented a human-in-the-loop (HIL) retraining pipeline, enabling continuous model improvement from user-reviewed corrections.
- Built an experimentation framework for prompt engineering, model selection, and evaluation, including **A/B testing** across LLM providers to optimize accuracy, latency, and cost.
- Created evaluation pipelines to measure extraction performance (precision, recall, F1, CER) against ground truth, enabling data-driven iteration and model selection.
- Built backend services in Python to support real-time and batch inference for document processing pipelines.

US Bank | Machine Learning Engineer | Seattle, WA

Jan 2025 – April 2025

- Migrated onprem MySQL database to SQL DB hosted in Azure, updating machine learning models and backend server in Python to call the new database.
- Configured permissions in Azure Cloud and Kubernetes to send 10TB daily of financial records data through Kafka service and persist data in Azure SQL database.
- Updated CI/CD pipeline in GitLab to run 5 additional test cases for new connection to Azure database.

HundredX | AI and Machine Learning Engineer | Seattle, WA

Feb 2022 – May 2024

- Led cross-functional initiative to transform customer feedback into actionable insights using NLP and LLM-based systems, retaining **\$120K+ in enterprise contracts**.
- Built and fine-tuned NLP models (BERT, Transformers) for topic modeling and sentiment analysis, achieving **93% accuracy**.
- Leveraged LLMs (OpenAI API) for semantic understanding and classification of customer comments at scale, processing **10K+ comments daily** via Airflow, Docker, and Kubernetes, delivering insights to BigQuery.
- Improved sentiment model performance by **up to 33%** through model experimentation and evaluation.
- Developed Tableau dashboards and SQL pipelines to surface trends and insights to clients.
- Established model evaluation standards and experimentation workflows for ML systems.

Amazon | Software Engineer | Seattle, WA

July 2019 – Dec 2021

- Deployed ML model to reduce errors in manual data annotation by 10% utilizing React, Java, and AWS SageMaker.
- Designed and developed data pipelines to aggregate customer behavior data inside Amazon Go stores for the business intelligence team, helping land contracts valued at over **\$5M** using Java, AWS (DynamoDB, Kinesis, S3).
- Researched machine learning models and debugged errors and discrepancies to reduce checkout errors by 10% and 10 hours of manual work weekly; models ran in Python and Java with AWS EC2 distributed system architecture.
- Designed end-to-end test suite to ensure quality and implemented a continuous integration (CI) pipeline, saving 3 hours of manual engineer labor weekly using Java, JUnit, and AWS (CloudFormation, Monitoring, and Alarming).

Spectrum Effect | Software Engineer | Monterrey, NL, Mexico

Jan 2017 – July 2019

- Implemented machine learning algorithms to identify network issues for 3 major US cellular carriers, leading to contracts over **\$500K** using Python, Pandas, SciPy, Docker, Kubernetes, Neo4j, and Kafka.
- Designed data pipelines to train classification models with **96% accuracy** using Kafka, Redis, Kubernetes, and Python.
- Implemented DevOps practices to maintain systems parsing 3TB of binary data from clients every 15 minutes using multiprocessing with Python, Docker, and Kafka.
- Maintained integrity of production environments by implementing software reliability practices such as rollbacks using version control (Git, Bitbucket), CI/CD practices (TeamCity), and knowledge sharing (Confluence).